

Parallel Computation

02-04-2020

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.0 --

## v ggplot2 3.2.1    v purrr  0.3.3
## v tibble  2.1.3    v dplyr  0.8.3
## v tidyr   1.0.0    v stringr 1.4.0
## v readr   1.3.1    v forcats 0.4.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

Parallel Computation

Implicit Parallel

- BLAS (Basic Linear Algebra Subroutines)
 - CRAN R shippings with a version of single threaded BLAS library.
 - Microsoft R Open ships with Intel MKL (Win/Linux) / Accelerate ML (macOS) BLAS libraries.
 - on macOS, R could be configured to use the optimized BLAS from Apple's Accelerate framework
 - We could only install R with different BLAS libraries such as openblas or ATLAS

Embarrassingly Parallel

Also called perfectly parallel, delightfully parallel or pleasingly parallel.

An embarrassingly parallel task can be considered a trivial case - little or no manipulation is needed to separate the problem into a number of parallel tasks.

A bit deroute first - revisit some of our old friends `map` and `map_*` in `purrr`.

```
1:4 %>% map(function(x) x^2)
```

```
## [[1]]
## [1] 1
##
## [[2]]
## [1] 4
```

```
##
## [[3]]
## [1] 9
##
## [[4]]
## [1] 16
```

```
1:4 %>% map_dbl(function(x) x^2)
```

```
## [1] 1 4 9 16
```

```
1:4 %>% map_dbl(~ .^2)
```

```
## [1] 1 4 9 16
```

These are the base R equivalence.

```
1:4 %>% lapply(function(x) x^2)
```

```
## [[1]]
## [1] 1
##
## [[2]]
## [1] 4
##
## [[3]]
## [1] 9
##
## [[4]]
## [1] 16
```

```
1:4 %>% sapply(function(x) x^2)
```

```
## [1] 1 4 9 16
```

Suppose we have a list of vectors and we want to operation some operation on each vector.

```
# it is a slow operation, imagine that in real applications, it could take a few minutes
slow_task <- function(x) {
  sum(x %o% x)
}
```

```
list_of_vectors <- replicate(10, list(rnorm(5000)))
list_of_vectors %>% glimpse()
```

```
## List of 10
## $ : num [1:5000] 0.775 0.639 0.107 0.533 -0.15 ...
## $ : num [1:5000] 0.313 0.879 -0.194 -2.059 1.447 ...
## $ : num [1:5000] 0.402 0.3066 0.5407 -0.5085 -0.0382 ...
## $ : num [1:5000] 0.802 -0.218 1.087 -0.222 0.774 ...
```

```
## $ : num [1:5000] -0.9784 1.1896 0.0203 0.3901 0.6679 ...
## $ : num [1:5000] 0.862 1.317 0.539 0.695 -0.761 ...
## $ : num [1:5000] -2.271 -0.991 0.563 1.537 0.515 ...
## $ : num [1:5000] -1.472 0.469 -1.147 1.536 -0.544 ...
## $ : num [1:5000] 0.662 0.318 0.734 -0.716 -1.115 ...
## $ : num [1:5000] -0.364 -0.215 0.259 0.668 -0.102 ...
```

```
list_of_vectors %>% map_dbl(slow_task)
```

However, these commands only run in a single process, it means, if the list is doubled, the time is also at least doubled.

```
system.time({
  list_of_vectors %>% map_dbl(slow_task)
})
```

```
##      user  system elapsed
##    2.468    1.141     3.683
```

```
# double the list
longer_list_of_vectors <- c(list_of_vectors, list_of_vectors)
system.time({
  longer_list_of_vectors %>% map_dbl(slow_task)
})
```

```
##      user  system elapsed
##    4.241    2.146     6.631
```

We are hoping to use multiple processes to speed up the job. The traditional way is to use the `parallel` package.

The package `parallel`

```
library(parallel)
```

Consider again the above `list_vector` example,

```
# the number of cores we have
detectCores()
```

```
## [1] 8
```

```
# it will create a socket cluster on my own computer
cl <- makeCluster(4)
parLapply(cl, list_of_vectors, slow_task)
```

```
## [[1]]
## [1] 6466.426
```

```
##
## [[2]]
## [1] 27662.34
##
## [[3]]
## [1] 2389.771
##
## [[4]]
## [1] 39.54569
##
## [[5]]
## [1] 817.6046
##
## [[6]]
## [1] 403.4123
##
## [[7]]
## [1] 10655.5
##
## [[8]]
## [1] 1159.628
##
## [[9]]
## [1] 3966.624
##
## [[10]]
## [1] 434.5246
```

```
# or if you want simplified result
parSapply(cl, list_of_vectors, slow_task)
```

```
## [1] 6466.42626 27662.34303 2389.77135 39.54569 817.60457 403.41228
## [7] 10655.50161 1159.62790 3966.62361 434.52465
```

```
# stop the cluster after use
stopCluster(cl)
```

Remark: you don't have to make and stop clusters for every operation, you could make a cluster in the very beginning of your script and close it at the very end.

Let's test the speed improvement

```
run_each <- function(x, fun, n_cores) {
  cl <- makeCluster(n_cores)
  result <- parLapply(cl, x, fun)
  stopCluster(cl)
  result
}

system.time(run_each(longer_list_of_vectors, slow_task, 2))
```

```
## user system elapsed
## 0.009 0.007 3.158
```

```
system.time(run_each(longer_list_of_vectors, slow_task, 8))
```

```
##      user  system elapsed  
## 0.033   0.041   6.759
```

Processing Chunk

The iterable is divided into chunks before sending the chunks to the workers. `Sys.getpid()` tells us the process id of a worker.

```
cl <- makeCluster(4)
```

```
parSapply(cl, 1:10, function(x) {  
  Sys.getpid()  
})
```

```
## [1] 8600 8600 8600 8617 8617 8629 8629 8641 8641 8641
```

```
parSapply(cl, 1:10, function(x) {  
  Sys.getpid()  
},  
chunk.size = 2  
)
```

```
## [1] 8600 8600 8617 8617 8629 8629 8641 8641 8600 8600
```

```
parSapply(cl, 1:10, function(x) {  
  Sys.getpid()  
},  
chunk.size = 1  
)
```

```
## [1] 8600 8617 8629 8641 8600 8617 8629 8641 8600 8617
```

```
stopCluster(cl)
```

Load balancing

`parLapply` pre-schedules the tasks to each work. It could be suboptimal when different tasks require different amount of time to complete.

```
cl <- makeCluster(4)
```

```
x <- c(3, 3, 1, 1, 1, 1, 1, 1)  
pause <- function(x) {  
  Sys.sleep(x)  
}  
  
system.time({  
  map(x, pause)  
})
```

```
##      user  system elapsed
##    0.005   0.000  12.016
```

```
system.time({
  parLapply(cl, x, pause)
})
```

```
##      user  system elapsed
##    0.003   0.000   6.010
```

```
system.time({
  parLapply(cl, x, pause, chunk.size = 2)
})
```

```
##      user  system elapsed
##    0.002   0.001   6.006
```

```
system.time({
  parLapply(cl, x, pause, chunk.size = 1)
})
```

```
##      user  system elapsed
##    0.003   0.000   4.009
```

Instead of prescheduling the tasks, a task could also be assigned to a free worker dynamically using `parLapplyLB`.

```
system.time({
  parLapplyLB(cl, x, pause)
})
```

```
##      user  system elapsed
##    0.006   0.001   3.007
```

Note that it only takes 3 seconds now.

```
stopCluster(cl)
```

Interact directly with the workers

We just saw an quick example on using `parLapply`. Let's try a few more things.

```
cl <- makeCluster(4)
```

We could run some arbitrary commands on each of the workers

```
clusterEvalQ(cl, {
  x <- rnorm(100)
  mean(x)
})
```

```
## [[1]]
## [1] 0.02233895
##
## [[2]]
## [1] -0.1053939
##
## [[3]]
## [1] -0.03107646
##
## [[4]]
## [1] 0.003768593
```

```
clusterEvalQ(cl, {
  Sys.getpid()
})
```

```
## [[1]]
## [1] 8719
##
## [[2]]
## [1] 8732
##
## [[3]]
## [1] 8744
##
## [[4]]
## [1] 8756
```

Global variables in master are not exported to the worker automatically (the same is true for `parLapply`)

```
y <- 3
clusterEvalQ(cl, {
  y + 1
})
```

```
## Error in checkForRemoteErrors(lapply(cl, recvResult)): 4 nodes produced errors; first error: object
```

`clusterExport` exports the global variables to each worker.

```
clusterExport(cl, "y")
clusterEvalQ(cl, {
  y + 1
})
```

```
## [[1]]
## [1] 4
##
## [[2]]
## [1] 4
##
## [[3]]
## [1] 4
```

```
##  
## [[4]]  
## [1] 4
```

Note: if you use `clusterExport` inside a function, you may want to specify the `envir` parameter.

```
doCalc <- function() {  
  z <- 5  
  # export z from the current scope  
  clusterExport(cl, "z", envir = environment())  
  clusterEvalQ(cl, {  
    z + 1  
  })  
}  
doCalc()
```

```
## [[1]]  
## [1] 6  
##  
## [[2]]  
## [1] 6  
##  
## [[3]]  
## [1] 6  
##  
## [[4]]  
## [1] 6
```

If you want to set a random seed, the following doesn't work because each work returns the same result.

```
# wrong  
set.seed(123)  
clusterEvalQ(cl, {  
  rnorm(5)  
})
```

```
## [[1]]  
## [1] -0.88326009 -2.16480736  1.77660023 -0.13725486  0.06906422  
##  
## [[2]]  
## [1] -0.8128060  0.4333262 -0.2760510 -1.1120264 -1.0612642  
##  
## [[3]]  
## [1] -0.03124709  2.15854565 -2.47045119 -0.01678581 -0.41465096  
##  
## [[4]]  
## [1]  0.4309572 -0.2502550  1.6120490  0.8735261 -1.6402650
```

```
# wrong again  
clusterEvalQ(cl, {  
  set.seed(123)  
  rnorm(5)  
})
```



```
## [[1]]
## [1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774
##
## [[2]]
## [1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774
##
## [[3]]
## [1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774
##
## [[4]]
## [1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774
```

```
clusterSetRNGStream(cl, 123)
clusterEvalQ(cl, {
  rnorm(5)
})
```

```
## [[1]]
## [1] -0.9685927  0.7061091  1.4890213 -1.8150926  0.3304096
##
## [[2]]
## [1] -0.4094454  0.8909694 -0.8653704  1.4642711  1.2674845
##
## [[3]]
## [1] -0.48906078  0.43304237 -0.03195349  0.14670372 -1.75239095
##
## [[4]]
## [1] -1.0388664  1.5745125  0.7470820  0.6718720  0.2691436
```

```
# do not forget to close the cluster
stopCluster(cl)
```

Another popular package: foreach

As we have seen a bit earlier, we might need to use `clusterExport` to export certain global variables. It would be a bit cumbersome. To reduce the extra steps, we could consider `foreach` and `doParallel`. They will send all the globals to the workers before running the tasks. (However, I don't personally recommend `foreach` though it is (was!?) popular, I prefer `furrr` in the below)

```
library(foreach)
```

```
##
## Attaching package: 'foreach'

## The following objects are masked from 'package:purrr':
##
##   accumulate, when
```

```
library(doParallel)
```

```
## Loading required package: iterators
```

```

cl <- makeCluster(4)
registerDoParallel(cl)
m <- matrix(rnorm(9), 3, 3)
# matrix m is sent to the workers implicitly
# in fact, all globals are sent to the workers in default
foreach(i = seq_len(nrow(m)), .combine = c) %dopar% {
  sum(m[i, ])
}

```

```
## [1] -0.02905105 -1.36595099  2.58692045
```

```
stopCluster(cl)
```

However, it is not quite “functional” as `map`, `lapply` and `parLapply`.

map or lapply like syntax

mclapply from parallel (unix / macOS only)

Remark: `mclapply` relies on forking, it means that it doesn’t work on Windows. We will discuss a cross platform approach.)

```

# in default, `mclapply` uses 2 cores
system.time({
  mclapply(
    c(2, 2, 2, 2),
    function(x) Sys.sleep(x)
  )
})

```

```
##      user  system elapsed
##  0.005   0.012   4.015
```

```

system.time({
  mclapply(
    c(3, 3, 1, 1, 1, 1, 1, 1),
    function(x) Sys.sleep(x),
    mc.preschedule = FALSE, # set FALSE to enable load balancing
    mc.cores = 4
  )
})

```

```
##      user  system elapsed
##  0.020   0.034   3.028
```

package furrr

`furrr` provides functions which are very similar to those in `purrr`. One nice thing about `furrr` is that it doesn’t send absolutely every global variables to the workers as `foreach` and `mclapply`. It does lexical analysis to deduce what objects are needed and only transfer them to the workers. It also loads libraries automatically in the workers.

```
library(furrr)
```

```
## Loading required package: future
```

```
# to use 4 workers, `plan(multiprocess)` will use all the available workers
plan(multiprocess, workers = 4)
system.time({
  future_map(
    c(2, 2, 2, 2),
    ~ Sys.sleep(.)
  )
})
```

```
##      user  system elapsed
##    0.089   0.003   2.388
```

future_map has a family of type specific functions. For example,

```
future_map_dbl(list(1:10, 11:20, 21:30, 31:41), ~ sum(.))
```

```
## [1] 55 155 255 396
```

load balanacing in future_map

```
# without load balanacing
system.time({
  future_map(
    c(3, 3, 1, 1, 1, 1, 1, 1),
    ~ Sys.sleep(.)
  )
})
```

```
##      user  system elapsed
##    0.111   0.005   6.046
```

```
# with load balanacing
system.time({
  future_map(
    c(3, 3, 1, 1, 1, 1, 1, 1),
    ~ Sys.sleep(.),
    .options = future_options(scheduling = FALSE)
  )
})
```

```
##      user  system elapsed
##    0.171   0.009   4.008
```

The honorable mention multidyplr

The package multidyplr is tidyverse solution to multiprocess data management. However, it is not released to CRAN yet, it means we need to install it from its github repo by using the package `remotes`.

```
remotes::install_github("tidyverse/multidplyr")
```

Since it is still under development, use it with caution!

```
library(multidplyr)
```

There are two ways to get data to the workers in cluster:

- `partition()` a data frame that already loaded in the interactive process.
- Load a different subset of the data in each worker.

Use `partition()` to send data to workers

```
library(nycflights13)

cluster <- new_cluster(4)

flights %>%
  group_by(dest, origin) %>%
  partition(cluster) %>%
  summarize(air_time = mean(air_time, na.rm = TRUE))
```

```
## Source: party_df [224 x 3]
## Groups: dest
## Shards: 4 [43--67 rows]
##
##   dest  origin air_time
##   <chr> <chr>    <dbl>
## 1 ABQ   JFK      249.
## 2 ATL   JFK      112.
## 3 AVL   EWR       89.8
## 4 AVL   LGA       92.2
## 5 BDL   EWR       25.5
## 6 BHM   JFK       117
## # ... with 218 more rows
```

Load data in each worker

```
cluster <- new_cluster(4)

cluster %>% cluster_library(c("tidyverse", "nycflights13"))

cluster %>%
  cluster_assign_partition(destination = unique(flights$dest))

# let's check the `destination` variable in the workers
cluster %>%
  cluster_call(destination)
```

```
## [[1]]
## [1] "IAH" "MIA" "BQN" "ATL" "ORD" "FLL" "IAD" "MCO" "PBI" "TPA" "LAX" "SFO"
```

```
## [13] "DFW" "BOS" "LAS" "MSP" "DTW" "RSW" "SJU" "PHX" "BWI" "CLT" "BUF" "DEN"
## [25] "SNA" "MSY" "SLC"
##
## [[2]]
## [1] "XNA" "MKE" "SEA" "ROC" "SYR" "SRQ" "RDU" "CMH" "JAX" "CHS" "MEM" "PIT"
## [13] "SAN" "DCA" "CLE" "STL" "MYR" "JAC" "MDW" "HNL" "BNA" "AUS" "BTV" "PHL"
## [25] "STT" "EGE"
##
## [[3]]
## [1] "AVL" "PWM" "IND" "SAV" "CAK" "HOU" "LGB" "DAY" "ALB" "BDL" "MHT" "MSN"
## [13] "GSO" "CVG" "BUR" "RIC" "GSP" "GRR" "MCI" "ORF" "SAT" "SDF" "PDX" "SJC"
## [25] "OMA" "CRW"
##
## [[4]]
## [1] "OAK" "SMF" "TUL" "TYS" "OKC" "PVD" "DSM" "PSE" "BHM" "CAE" "HDN" "BZN"
## [13] "MTJ" "EYW" "PSP" "ACK" "BGR" "ABQ" "ILM" "MVY" "SBN" "LEX" "CHO" "TVC"
## [25] "ANC" "LGA"
```

```
cluster %>% cluster_send({
  df <- flights %>% filter(dest %in% destination)
})

cluster %>%
  party_df("df") %>%
  group_by(dest, origin) %>%
  summarize(air_time = mean(air_time, na.rm = TRUE)) %>%
  collect()
```

```
## # A tibble: 224 x 3
## # Groups:   dest [105]
##   dest origin air_time
##   <chr> <chr>     <dbl>
## 1 ATL   EWR        112.
## 2 ATL   JFK        112.
## 3 ATL   LGA        114.
## 4 BOS   EWR         40.3
## 5 BOS   JFK         38.5
## 6 BOS   LGA         37.9
## 7 BQN   EWR        196.
## 8 BQN   JFK        194.
## 9 BUF   EWR         51.2
## 10 BUF  JFK         57.1
## # ... with 214 more rows
```

Reference:

- R Programming for Data Science <https://bookdown.org/rdpeng/rprogdatascience/parallel-computation.html>